

POMDPs

POMDPs

- We've been living a lie:

```
s = observe(env)
```

Types of Uncertainty


Alleatory



Epistemic (Static)



Epistemic (Dynamic)




Interaction



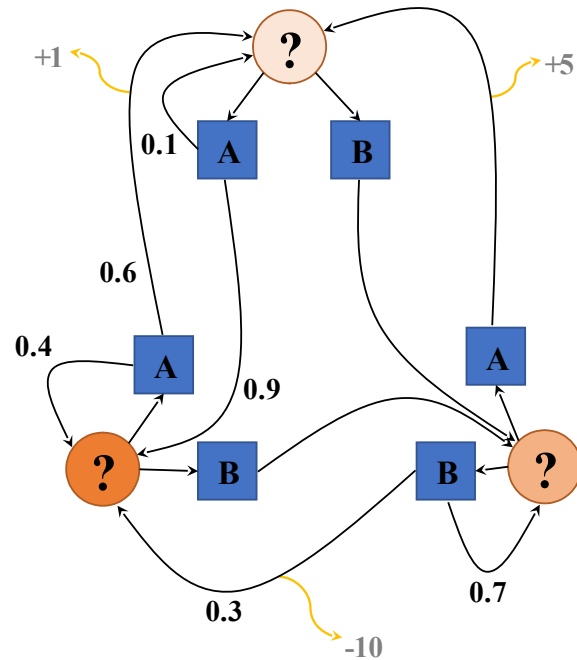
Markov Decision Process (MDP)

- \mathcal{S} - State space
- $T : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$ - Transition probability distribution
- \mathcal{A} - Action space
- $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ - Reward



Alleatory

Partially Observable Markov Decision Process (POMDP)



- \mathcal{S} - State space
- $T : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$ - Transition probability distribution
- \mathcal{A} - Action space
- $R : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$ - Reward
- \mathcal{O} - Observation space
- $Z : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \times \mathcal{O} \rightarrow \mathbb{R}$ - Observation probability distribution

Alleatory

Epistemic (Static)

Epistemic (Dynamic)

Tiger POMDP Definition

Hidden Markov Models and Beliefs

Bayesian Belief Updates

Filtering Loop

Tiger Example

Recap